

Nicholas Yegon

Matrikel Number: 1748461

Exercise 04

1. Implement Vanilla Gradient Descent (Not parallel)

Implement the algorithm from slide 8 As function use $f(x) := x^2$ with $\nabla f(x) := 2x$ Run the algorithm for 20 iterations with learning rate $\eta := 0.1$ and starting point 3.5 Plot the function in the interval $[-4, 4]$ and overlay the point that x visits in the 20 epochs The result should look like slide 9

```
import numpy as np
import matplotlib.pyplot as plt

def square_fn(x):
    return x**2

def gradient(x):
    return 2*x

def gradient_descent(square_fn, gradient, x_init, learning_rate, epsilon,
max_iterations):
    x = x_init
    x_values = [x]

    for i in range(max_iterations):
        gradient_x = gradient(x)
        x_next = x - learning_rate * gradient_x

        if np.abs(x_next - x) < epsilon:
            break

        x = x_next
        x_values.append(x)


    return x, x_values
```

```
learning_rate = 0.1
x_init = 3.5
num_iterations = 20
epsilon = 1e-6

x_optimal, x_values = gradient_descent(
    square_fn, gradient, x_init, learning_rate, epsilon,
    max_iterations=num_iterations)

x_axis = np.linspace(-4, 4, 100)
y_axis = square_fn(x_axis)
plt.plot(x_axis, y_axis)
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Gradient Descent')
plt.grid(True)

x_values_vals = np.array(x_values)
y_history_vals = square_fn(x_values_vals)
plt.scatter(x_values_vals, y_history_vals, c='red')
plt.legend()
plt.show()
```



2. Ramp Up

Generate data with the following code:

```
X = np . arange ( 0 , 1 , 0.01)
Y = X + np . random . normal ( 0 , 0.2 , len ( X ) )
```

Write a python function `y_hat(x, theta1, theta2)` that represents your model. Plot the data with `plt.scatter` and plot the models predictions with `plt.plot` using $\theta_1 = 1.0$ and $\theta_2 = 0.0$. (These are the optimal parameters)

```
import matplotlib.pyplot as plt
import numpy as np

def y_hat(x, theta1, theta2):
    return theta1 * x + theta2


# Generate data
X = np.arange(0, 1, 0.01)
```

```
Y = X + np.random.normal(0, 0.2, len(X))

# Plot the data
plt.scatter(X, Y)

# Plot the model's predictions
theta1 = 1.0
theta2 = 0.0
predictions = y_hat(X, theta1, theta2)
plt.plot(X, predictions, color='red')

plt.xlabel('x')
plt.ylabel('y')
plt.title('Model Predictions')
plt.legend()
plt.show()
```



Implement the algorithm from slide 13(non-parallel)

Initialize your parameters to $\theta_1 = -0.5$ and $\theta_2 = 0.2$, set the learning rate $\eta = 0.01$ and run for $E_{max} = 5$ iterations.

```
import numpy as np
import matplotlib.pyplot as plt

def y_hat(x, theta1, theta2):
    return theta1*x+theta2

#generate data
X = np.arange(0, 1, 0.01)
Y = X + np.random.normal(0, 0.2, len(X))

# Initialize parameters
theta1 = -0.5
theta2 = 0.2
```


```
# set learning rate and max iterations
learning_rate = 0.01
Emax = 5

# perform gradient descent
for i in range(Emax):

    # compute gradients
    grad_theta1 = np.mean((y_hat(X, theta1, theta2)-Y)*X)
    grad_theta2 = np.mean(y_hat(X, theta1, theta2)-Y)

    #Update params
    theta1 -= learning_rate * grad_theta1
    theta2 -= learning_rate * grad_theta2

# plot the data
plt.scatter(X,Y)
# plot the final model's predictions
predictions = y_hat(X, theta1, theta2)
plt.plot(X, predictions, color='red', label='model predictions')
plt.xlabel('x')
plt.ylabel('y')
plt.title('Model Predictions')
plt.legend()
plt.show()
```



3. Distributed Gradient Descent

Implement a distributed version of gradient descent by following slide 15

Test your code with different numbers of ranks and experiment with the amount of iterations needed to approach the optimal parameters $\theta_1 = 1.0$ and $\theta_2 = 0.0$.

```
import numpy as np
import matplotlib.pyplot as plt
from mpi4py import MPI

def y_hat(x, theta1, theta2):
    return theta1 * x + theta2

# Generate data
X = np.arange(0, 1, 0.01)
Y = X + np.random.normal(0, 0.2, len(X))

# Set the desired optimal parameters
```

```

desiredParam1 = 1.0
desiredParam2 = 0.0

# Initialize MPI
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
ranksSize = comm.Get_size()

# Partition the data
samples = len(X) // ranksSize
start = rank * samples
end = start + samples
varX = X[start:end]
varY = Y[start:end]

# Initialize parameters
theta1 = -0.5
theta2 = 0.2

# Set learning rate and max iterations
learning_rate = 0.01
Emax = 1000 # Increase the max iterations

# Perform gradient descent
for i in range(Emax):

    # Compute gradients locally
    gradientTheta1 = np.mean((y_hat(varX, theta1, theta2) - varY) * varX)
    gradientTheta2 = np.mean(y_hat(varX, theta1, theta2) - varY)

    # Sum gradients across ranks
    gradientTheta1_sum = comm.allreduce(gradientTheta1, op=MPI.SUM)
    gradientTheta2_sum = comm.allreduce(gradientTheta2, op=MPI.SUM)

    # Update parameters
    theta1 -= learning_rate * gradientTheta1_sum
    theta2 -= learning_rate * gradientTheta2_sum

    # Share the updated parameters with ranks
    comm.Bcast([theta1, MPI.DOUBLE], root=0)

```



```

comm.Bcast([theta2, MPI.DOUBLE], root=0)

# Check convergence
if np.abs(theta1 - desiredParam1) < 1e-6 and np.abs(theta2 -
desiredParam2) < 1e-6:
    break

# Gather the final model parameters to rank 0
finalTheta1 = comm.gather(theta1, root=0)
finalTheta2 = comm.gather(theta2, root=0)

# Plot the data
if rank == 0:
    plt.scatter(X, Y)

# Plot the final model's predictions
if len(finalTheta1) > 0 and len(finalTheta2) > 0:
    lastTheta1 = finalTheta1[-1]
    lastTheta2 = finalTheta2[-1]
    predictions = y_hat(X, lastTheta1, lastTheta2)

```

Do the needed iterations change with the amount of workers change? Yes, it is possible the iterations can change with the number of workers

Repeat this experiment with different model initializations